



Stream Processing Fundamentals

Randy May, Industry Solutions Advocate
March 2023



What You'll Learn

- ◆ The Goal: learn enough to feel comfortable tackling stream processing problems with Hazelcast
- ◆ You'll build a solution that processes an event stream and reacts in real-time.
- ◆ First run it on your laptop then scale it by deploying to Viridian
- ◆ Those who complete the workshop and deploy a Viridian cluster will earn a badge




Prerequisites

A Laptop Set Up for Java Development Including ...

- ✦ Java 11+
- ✦ Maven
- ✦ Docker Desktop
- ✦ Git
- ✦ A GitHub Account


Agenda

Introduction to the Pipeline API	4:00 - 4:15
Hand On Exercise	4:15 - 5:15
Deploy to Viridian	5:15 - 5:30



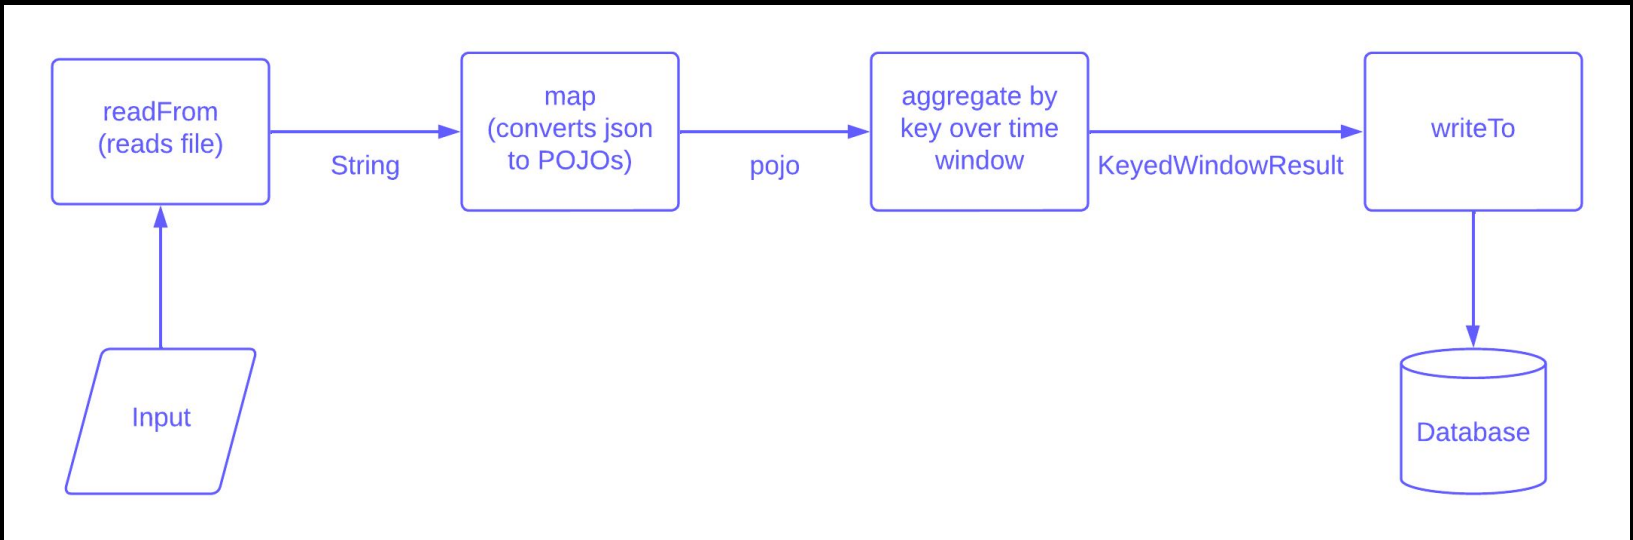


Let's Get Going ...




Hazelcast Pipelines

- ◆ A *Pipeline* is series of *tasklets* connected in a directed, acyclic graph.
- ◆ Events flow along the vertices of your Pipeline
- ◆ Events can be of any type
- ◆ Consider the simple Pipeline below



Pipeline Primitives



Hazelcast Pipelines

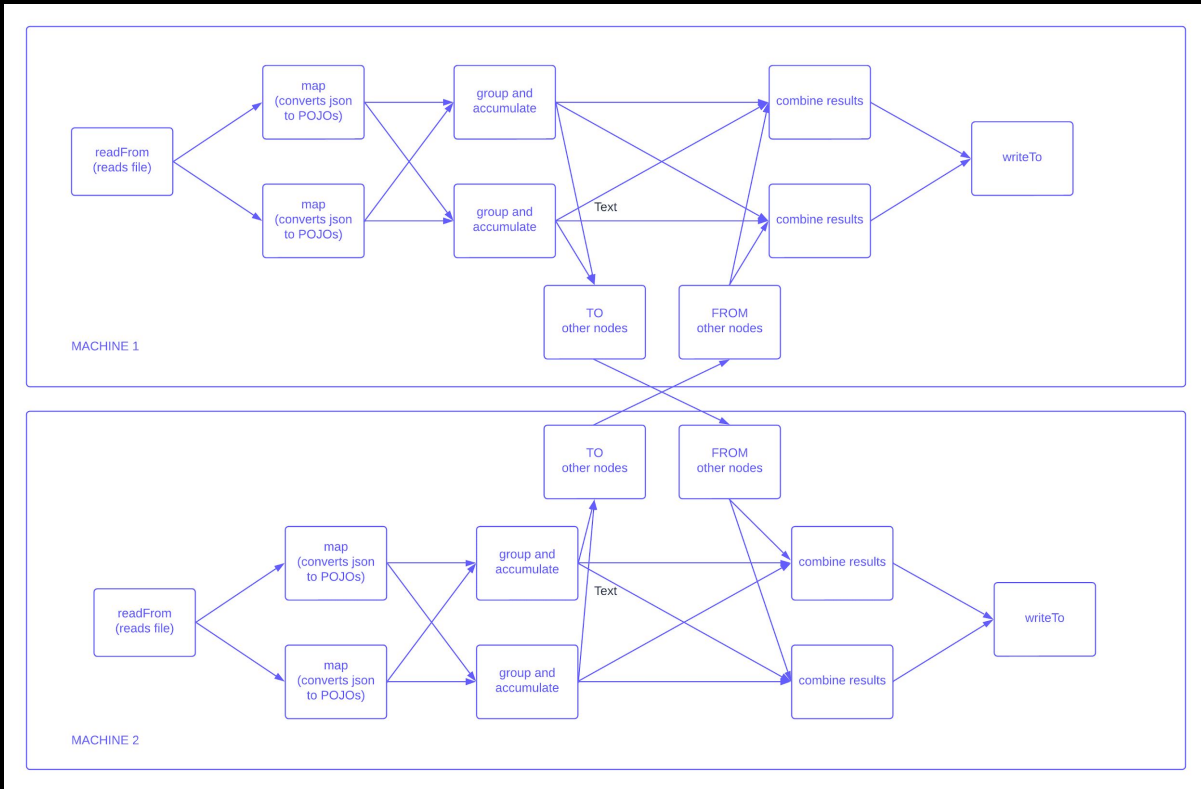
- ◆ You **define** the pipeline
- ◆ The Hazelcast platform **runs** it
- ◆ Your stream becomes ...
 - Parallelized
 - Location Aware
 - Fault Tolerant, Upgradeable and Restartable
- ◆ Scale by adding nodes

Hazelcast Pipelines

How the pipeline would be deployed on a cluster of 2 machines, each with 2 cores.

Much more detail is available here:

<https://docs.hazelcast.com/hazelcast/latest/architecture/distributed-computing>

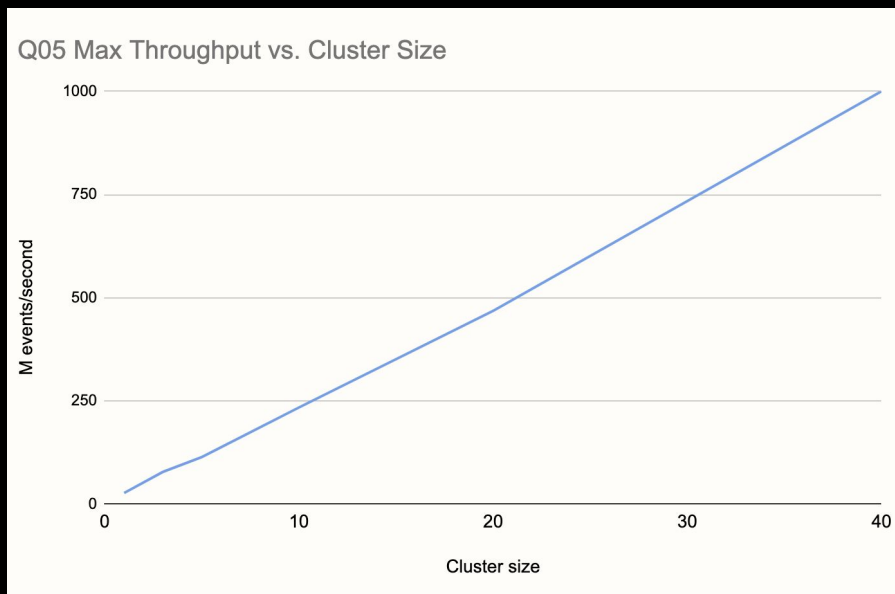



Why Learn The Pipeline API ?

1 Billion TPS with 99% under 30ms latency


- Keep-in-mind – the 100ms threshold from the Google Framework
- 45 nodes
- Linear scaling with predictable latency

[Link: Gigascale performance](#)





The Pipeline API



Example

Most methods return a typed `StreamStage`. The type parameters indicate what objects will flow along that vertex

From”

```
Pipeline pipeline = Pipeline.create();
```

```
StreamStage<Map.Entry<String, String>> requests =  
    pipeline.readFrom(  
        Sources.<String, Event>mapJournal(  
            REQUEST_MAP_NAME,  
            JournalInitialPosition.START_FROM_CURRENT))  
        .withTimestamps( event -> event.getEventTime());
```

Many methods take a functions as arguments. Often this is done with lambda expressions.

Example: adding labels using “map”

Input: (“abc”, 99.0)

Output: (“abc”, 99.0, “RED”)

Tuples are a useful helper



```
StreamStage<Tuple2<String,Integer>> temps;  
StreamStage<Tuple3<String,Integer,String>> labeledTemps =  
    temps.map(t -> Tuple3.tuple3(t.f0(),  
        t.f1(),  
        t.f1() > 90 ? "RED" : "GREEN"));
```

Lambdas again



Example: mapUsingService

Use when there is some external resource or connection that should not be re-created for each event (you'll need this)

```
StreamStage<MyPojo> stage;
```

ObjectMapper is created only once



```
ServiceFactory<?, ObjectMapper> objectMapperServiceFactory =  
    ServiceFactories.nonSharedService(ctx -> new ObjectMapper());
```

```
StreamStage<String> json =  
    stage.mapUsingService(objectMapperServiceFactory,  
        (mapper, pojo) -> mapper.writeValueAsString(pojo));
```



First arg is the “service”, second is the event to process

Example: windows and aggregation

For each machine, calculate the average temperature over a 10s “tumbling window”

```
StreamStage<MachineStatus> statusEvents;
```

Important! Specify the “group by” key


```
StreamStage<KeyedWindowResult<String, Double>> averageTemps =  
    statusEvents.groupingKey(s -> s.getSerialNum())  
        .window(WindowDefinition.tumbling(10000))  
        .aggregate(AggregateOperations.averagingLong(s -> s.getTemp()));
```

Window type and size

What and how to aggregate



The Lab

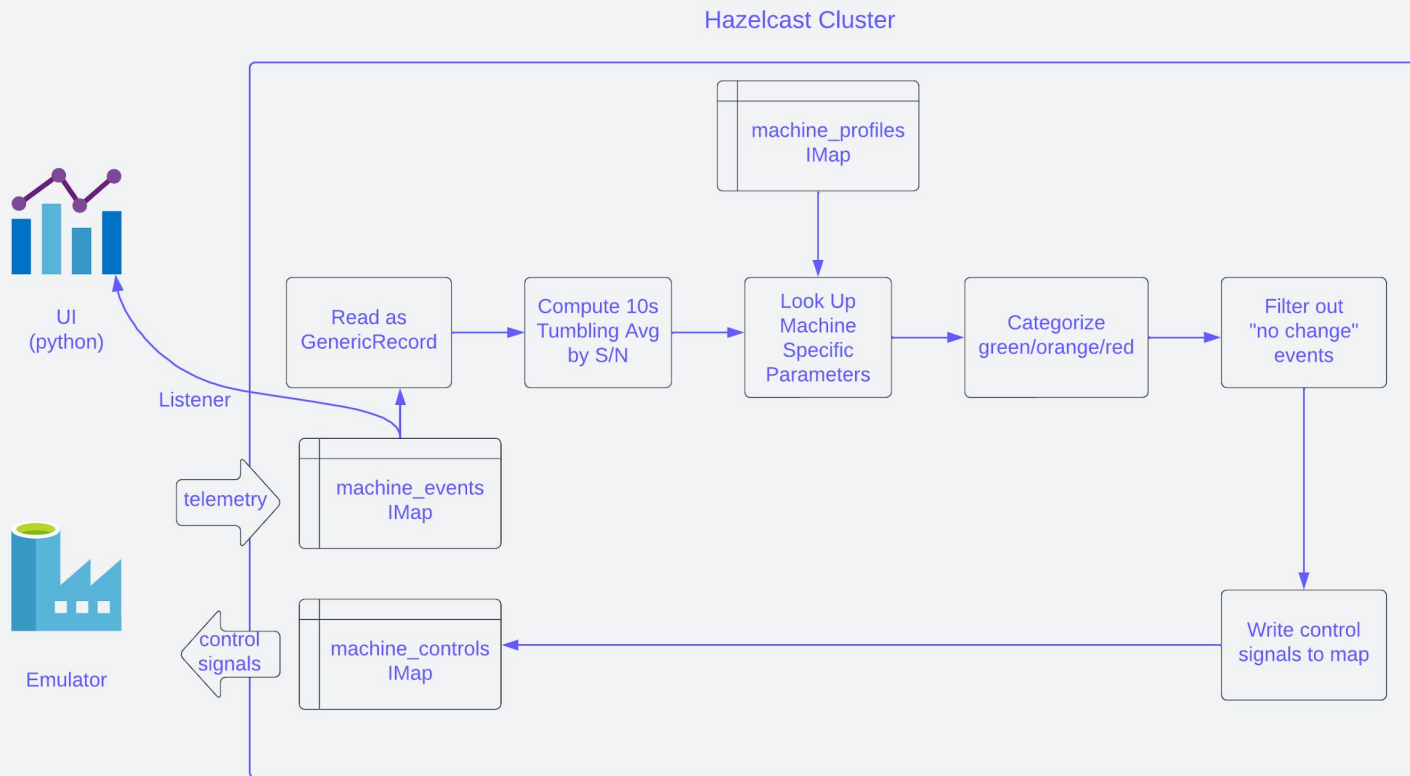


The Scenario

- ◆ Company operates 1000's of machines. Each publishes several data points each second. Measurements include things like bit temperature and RPM.
- ◆ Breakage is expensive. We want to go beyond maintenance schedules and monitor all of the information in real time.
- ◆ Each machine has its own parameters for acceptable bit temperature, which is stored in a “machine profile”
- ◆ If excessive bit temperatures are caught in time, breakage can be averted by immediately reducing the cutting speed



Pipeline for the Lab




Go Forth and Process Events!

- git clone
<https://github.com/hazelcast-guides/stream-processing-fundamentals>
- OR, if you've previously cloned the repo, **don't forget to “git pull”**
- Instructions and more detail are in the README
- Talk to the lab assistants!
- Feel free to reference the solution.



The Lab: 1 Hour





Deploy to Viridian

15 minutes

Note: if necessary, edit `cli/viridian_submitjob.sh` to submit the solution

Thank You !

Please help us make this workshop better by filling out the online survey

<https://www.surveymonkey.com/r/hzworkshop>



Quick Survey



How Do I Get Involved ?



slack.hazelcast.com



<https://github.com/hazelcast/hazelcast>



Networking

Food

Please Stick
Around !

Drinks

Nintendo
Switch
giveaway

Round Table
Discussion